



## Foundations

## Understanding the System

- Components

- Models

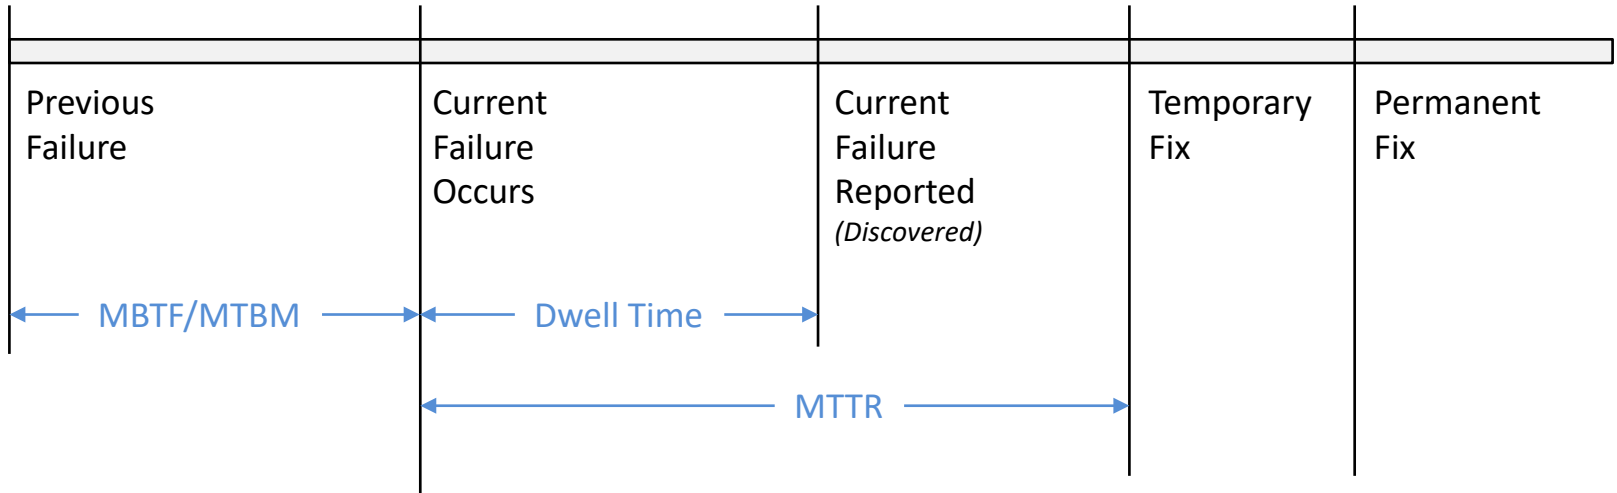
- System Knowledge

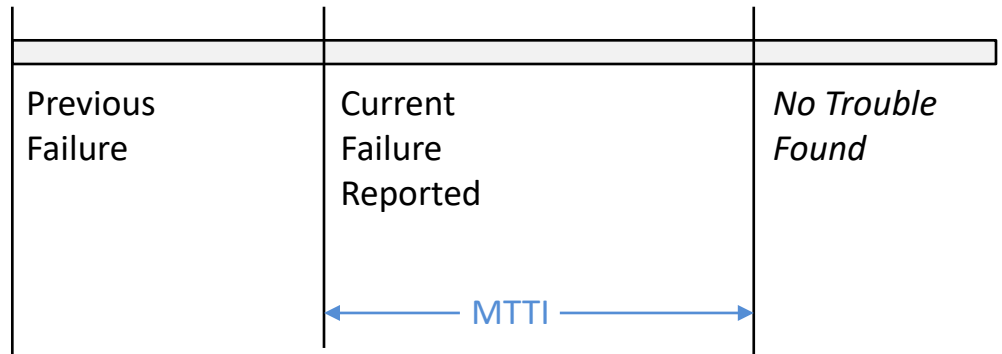
## Process

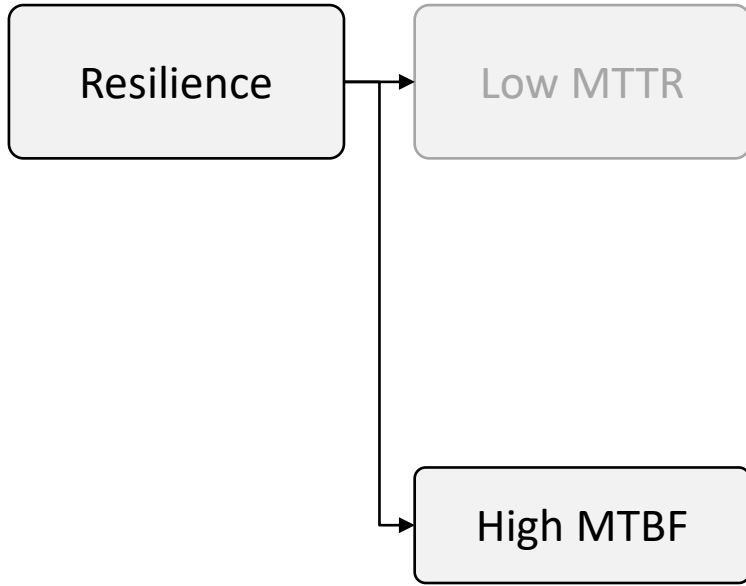
- The Half Split Method

## Some Examples









# Define “Down”

- A business critical application not operating at all?
- A business critical application operating slowly?
- The vice president cannot get to his email?
- An entire site is off line?
  - How large of a site?
- The president’s child downloaded a virus onto his personal computer?

MTBF 50k hours




MTBF 50k hours

$$\frac{1}{\frac{1}{MTBF_1} + \frac{1}{MTBF_2} + \frac{1}{MTBF_3} + \dots}$$

Combined MTBF: 16667 hours

MTBF 50k hours



MTBF 50k hours

Combined MTBF: 75000 hours  
 $\frac{1}{q} \left( 1 + \frac{1}{2} + \frac{1}{3} + \dots \right)$   
*Effectively increases MTBF by 50%*

$$\frac{1}{\frac{1}{MTBF_1} + \frac{1}{MTBF_2} + \frac{1}{MTBF_3} + \dots}$$

Combined MTBF: 18750 hours

2,5000 2ru 64x100g routers  
1,024 paths ToR to ToR

$$a_t = \frac{1}{\frac{1}{a_1} + \frac{1}{a_2} + \dots}$$

2,500 2ru 64x100g routers  
1,024 paths ToR to ToR

2,500 fabric routers  
MTBF 219,000 hours (25 years)

2,500 fabric routers x 64 ports each  
160,000 optical units  
MTBF 876,000 hours (100 years)

80,000 optical fibers  
160,000 optical connectors  
MTBF 876,000 hours (100 years)

$$a_t = \frac{1}{\frac{1}{a_1} + \frac{1}{a_2} + \dots}$$



2,5000 2ru 64x100g routers  
1,024 paths ToR to ToR

2,500 fabric routers  
MTBF 219,000 hours (25 years)

2,500 fabric routers x 64 ports each  
160,000 optical units  
MTBF 876,000 hours (100 years)

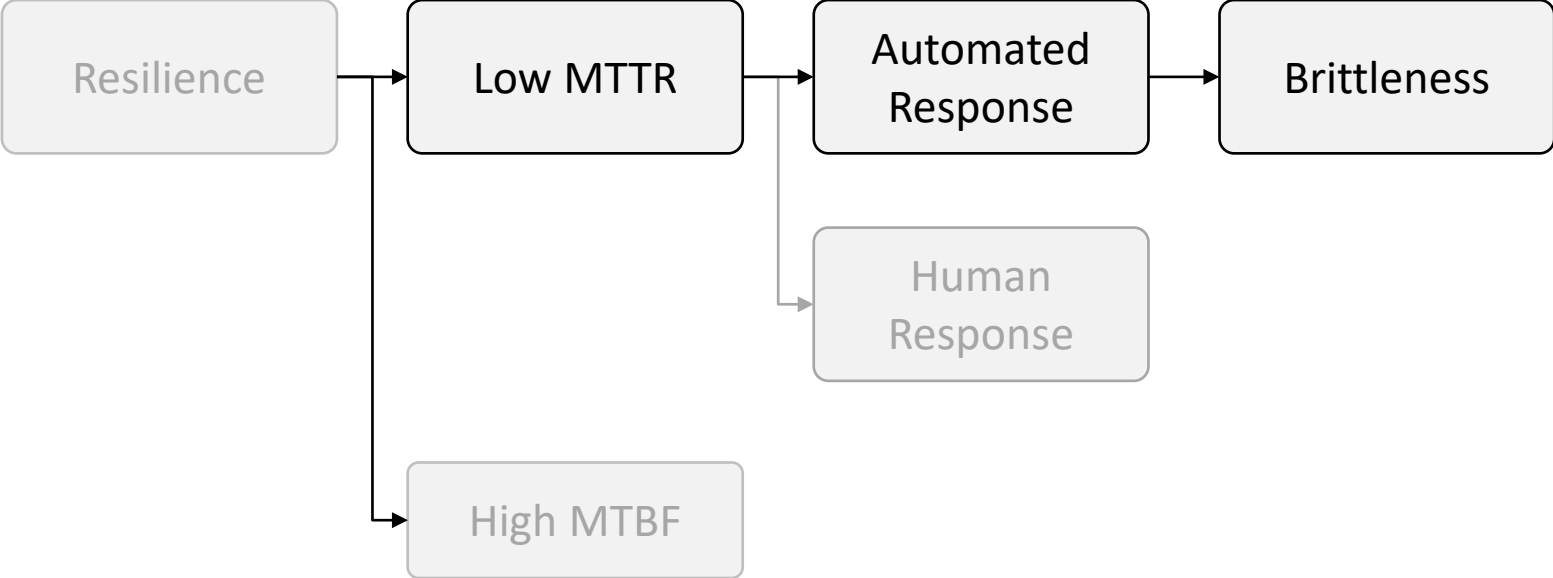
80,000 optical fibers  
160,000 optical connectors  
MTBF 876,000 hours (100 years)

$$a_t = \frac{1}{\frac{1}{a_1} + \frac{1}{a_2} + \dots}$$

1 router failure every 87.6 hours


1 optical failure every 5.5 hours

1 optical failure every 5.5 hours



# Automated Responses

- Automated responses can decrease MTTR
  - When they are well designed
  - When they function as intended
- Automated responses can increase MTTR
  - When they are not well designed
  - When they cause unintended consequences

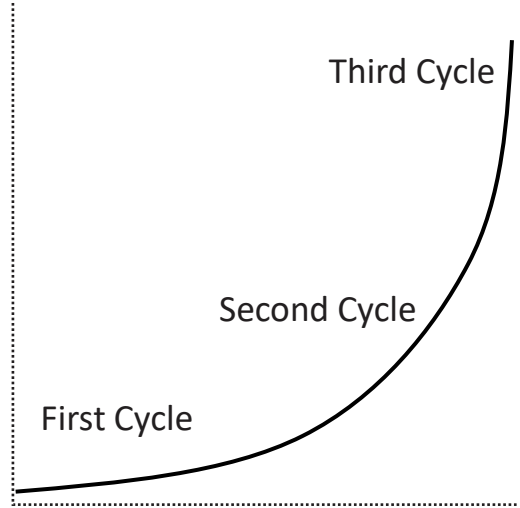
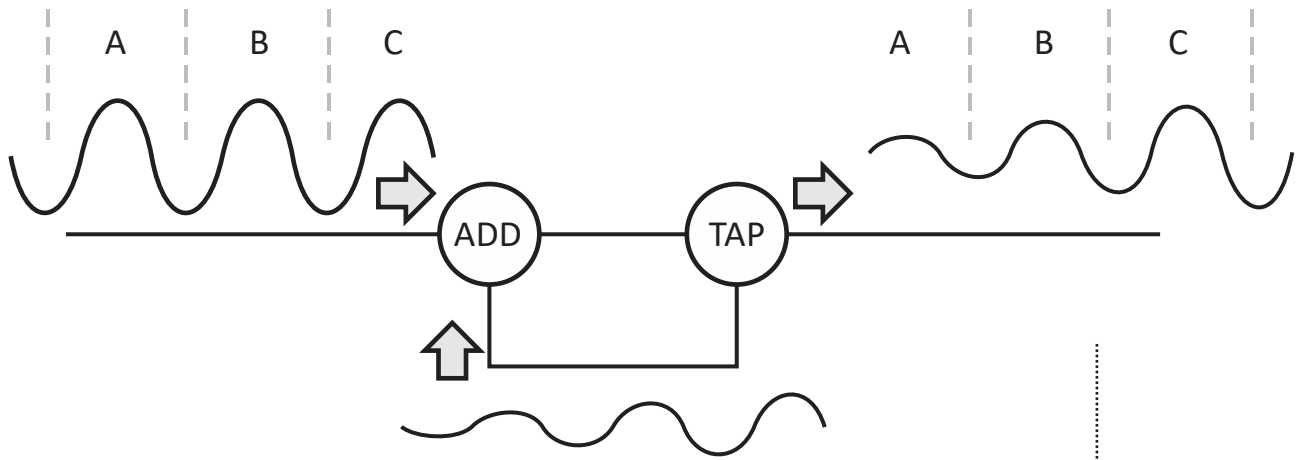



In a complex system, the cumulative effect of a large number of small optimizations is externally indistinguishable from a radical leap.

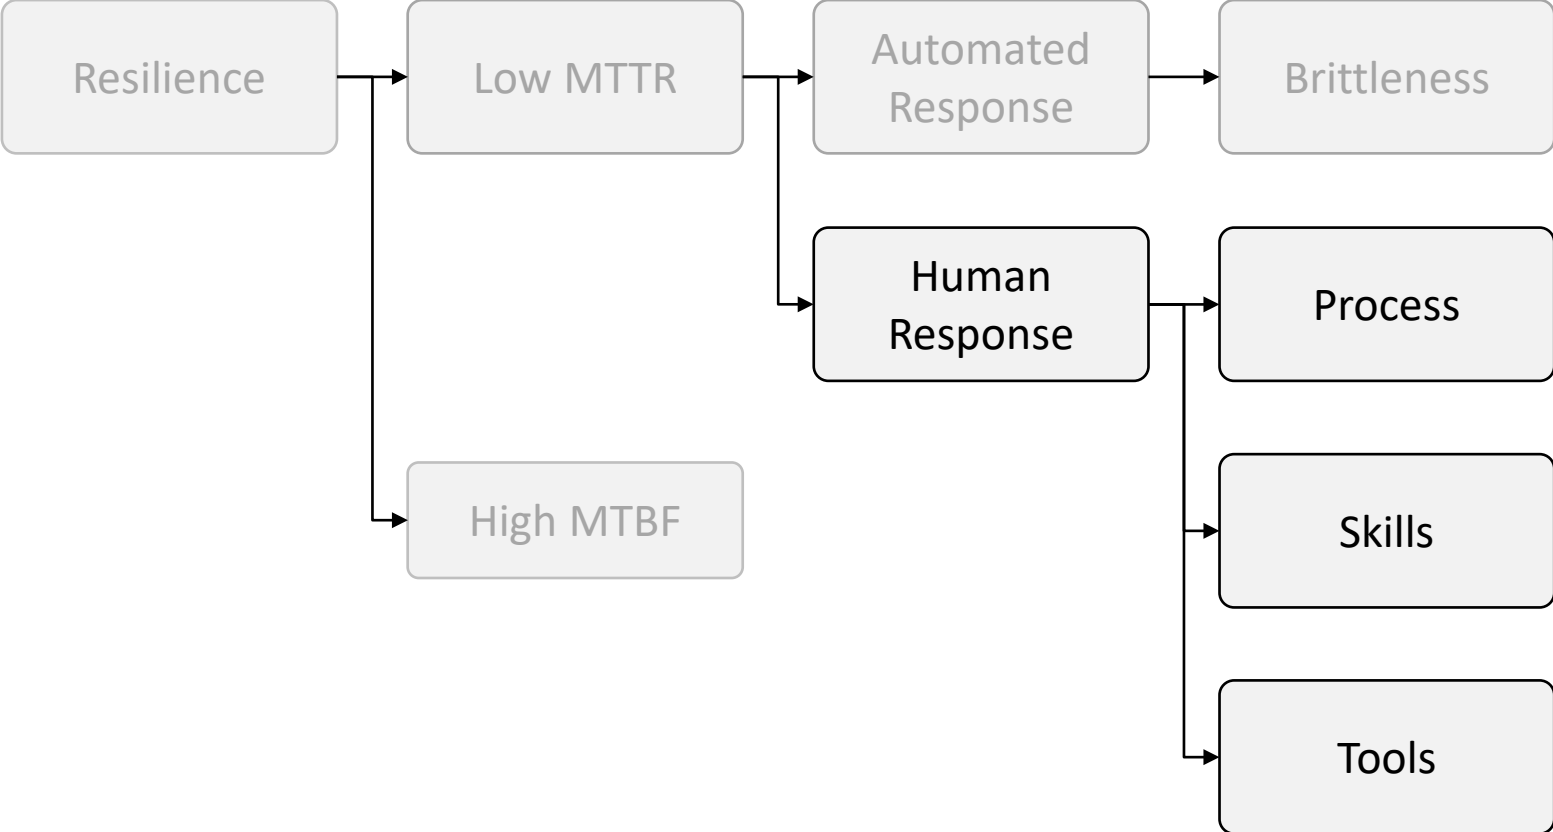
---

In a complex system, you can know your component in depth and the components around yours at a high level. The rest is rumor and pop psychology

*Keith's Law and the first corollary*







# Fixes

- Temp Fix
  - “We know this is going to break again quickly”
  - “We know this increases technical debt”
- Permanent Fix
  - “We know this won’t break anytime soon”
  - “This leaves technical debt where it was before the failure”
  - “This refactors the system and reduces technical debt”
- Technical debt
  - Mismatch between how the system really works and how you think it works





Components and Abstractions

Models


Measuring Things

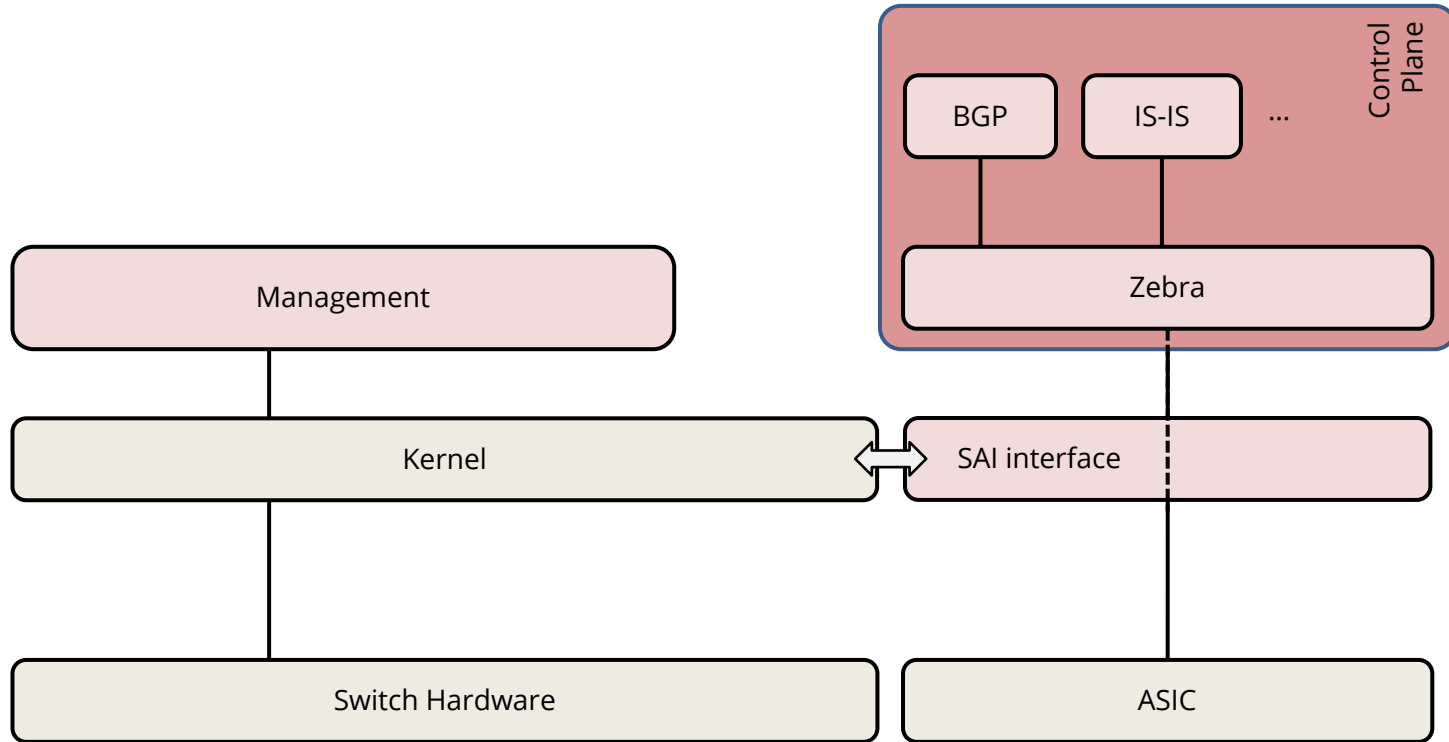
The Heisenberg Problem

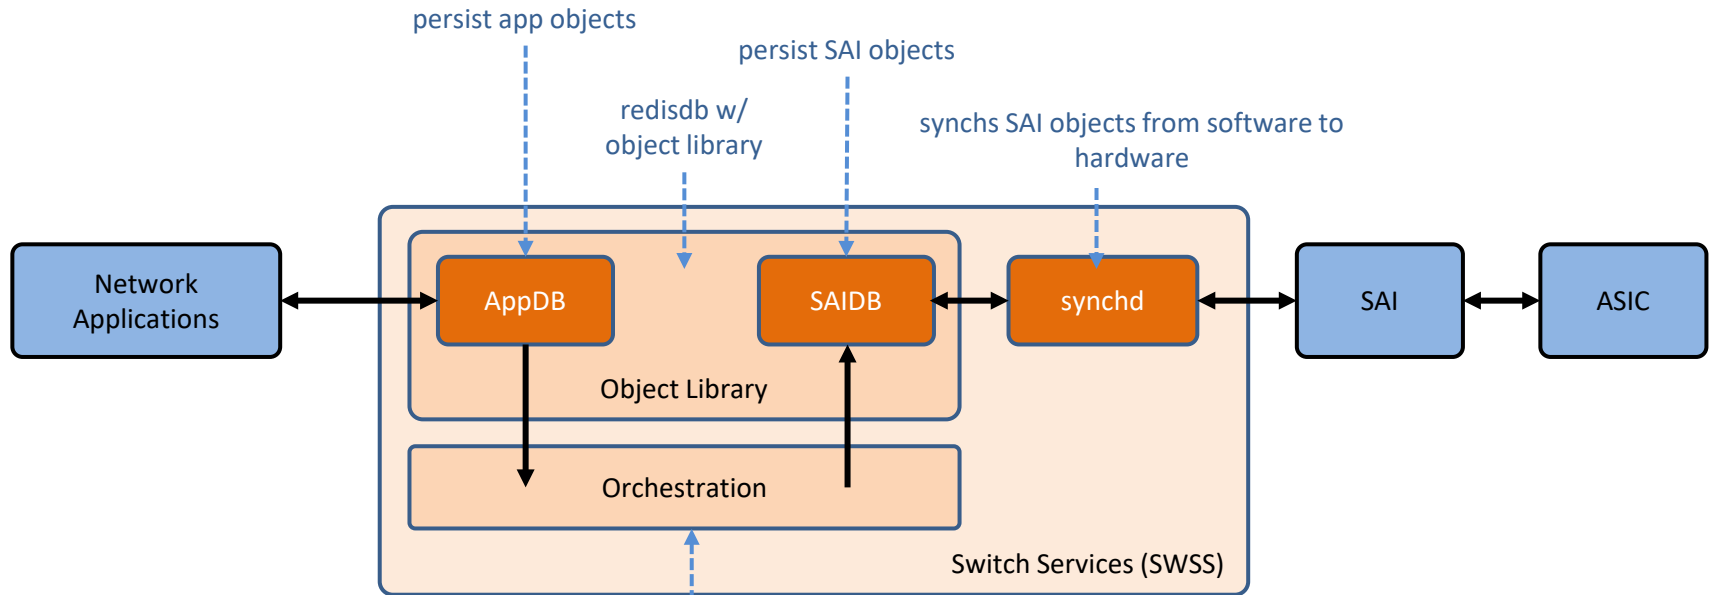
Manipulability Theory

# Understand the components

- You see a router, right?
- Is that all you see?

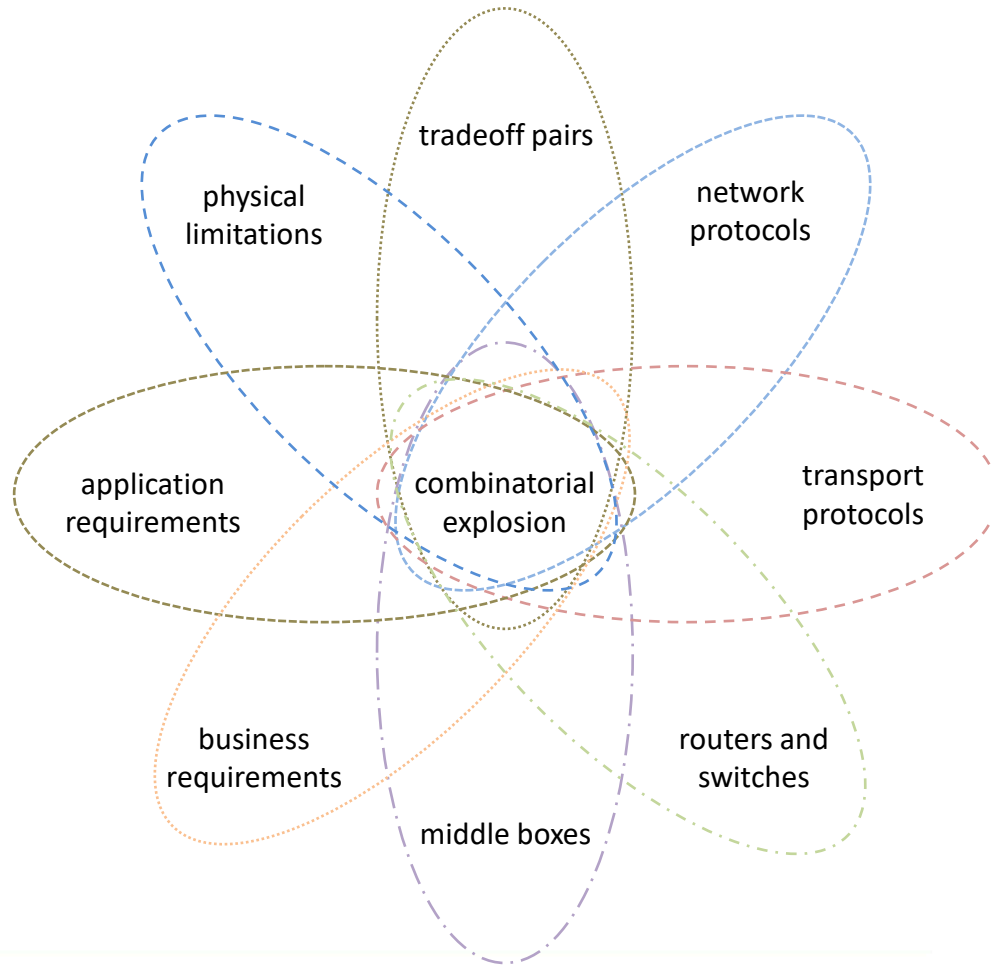






- translation between apps and SAI objects
- resolution of dependency and conflict





# Abstract!

# Abstract!

---

*Abstraction hide details*





# Abstract!

---


*Abstraction hide details*

---

*Abstraction leak*



*each colored line in the graph represents a single flow*



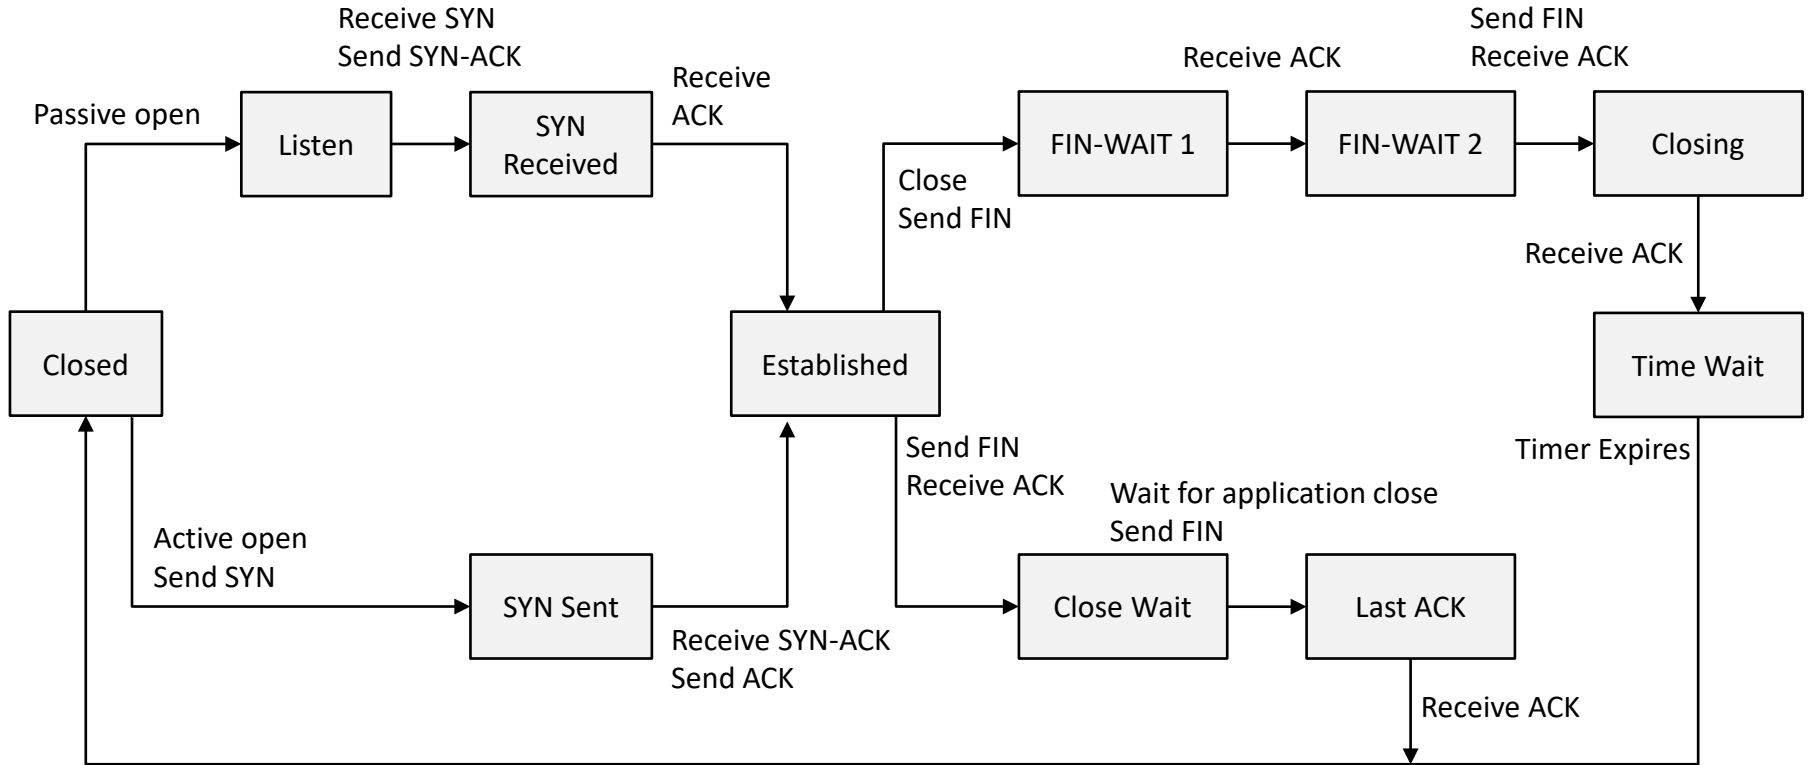
*each colored line in the graph represents a single flow*

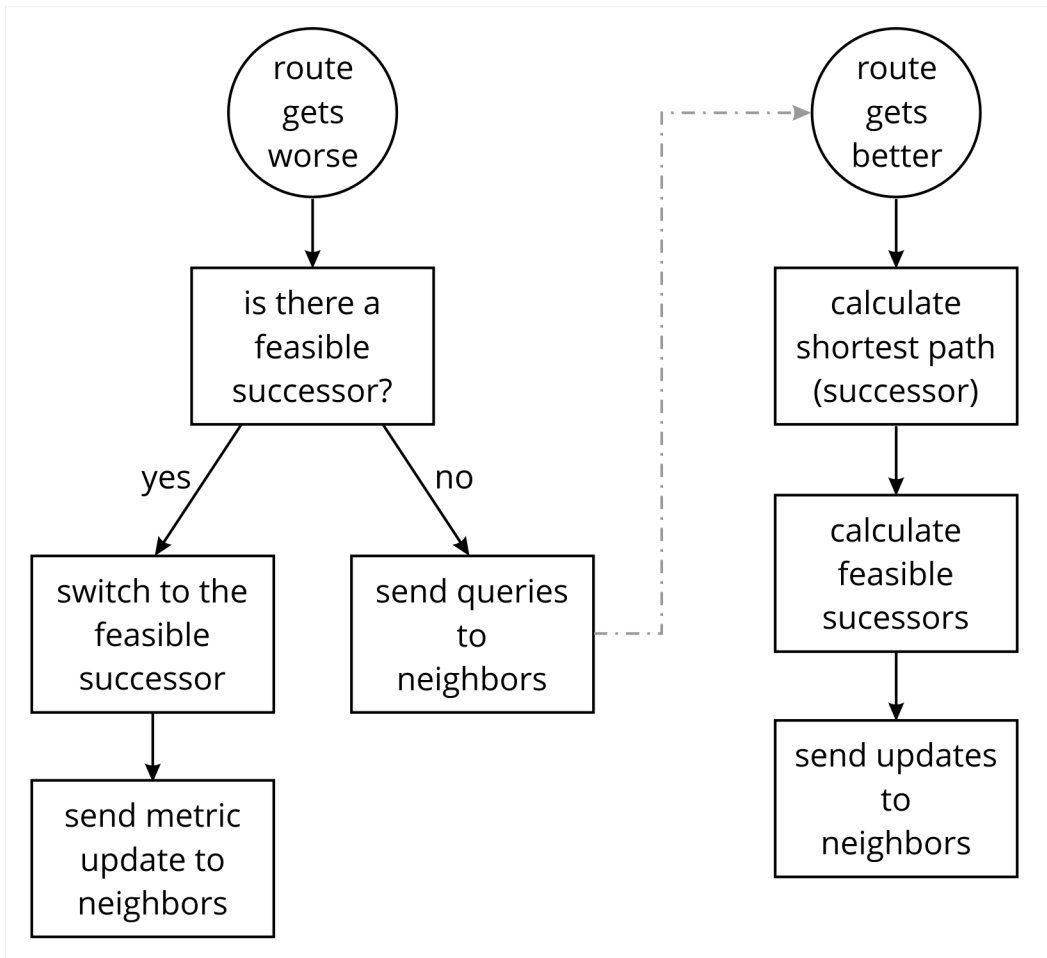
# Models

- Models are a form of abstraction
- What models
  - What does this do?
  - What happens when ... ?

# Models

- Models are a form of abstraction
- What models
  - What does this do?
  - What happens when ... ?
- How/why models
  - How does this happen?
  - Why does that happen this way?
- Some examples will help make sense of the differences





BGP  
Boundary  
Gateway  
Protocol

block message  
block length 2 bytes  
version number 1 byte  
block type 2 bytes (asymmetric)  
holddown timer 2 bytes (minutes)

types:  
open - 1  
update - 2  
notification - 4  
keepalive - 8  
version is currently 1

open:  
my AS # 2 byte  
link type 1 byte  
up - 1  
down - 2  
internal - 4 (not used in update structure field)  
H-link - 8  
auth type code 1 byte  
0 - none  
authentication variable

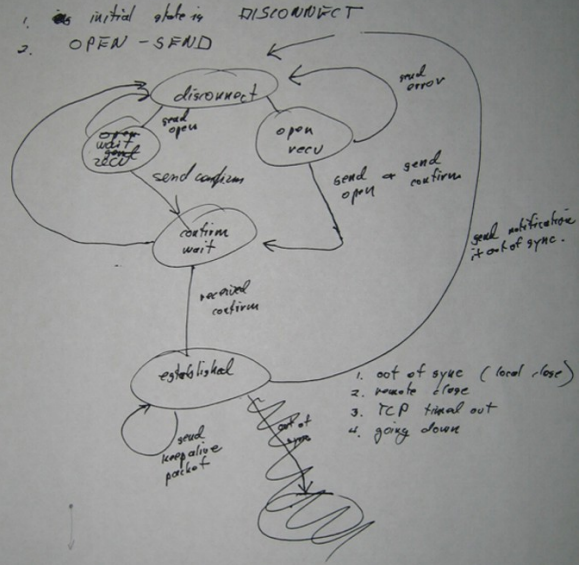
update:  
network # 4 bytes  
first hop gateway 4 bytes  
metric 2 bytes  
count of AS 1 byte  
direction 1 byte  
AS # 2 bytes  
repeat structure according to block length } repeat "route" times

notification:  
error code 2 bytes  
data variable

operation

- link type error in open  
- my view of correct link type (1 byte)
- unknown auth type code  
- no data
- authentication failure (no data)
- update error - data is block in error  
~~routing loop in update~~  
~~two phase error in update~~  
data is subcode (2 byte) followed by update block in question (1 network only)  
subcodes -  
1. invalid network field  
2. invalid first hop gw  
3. invalid direction code  
4. invalid AS  
5. routing loop  
6. two-phase error
- connection out of sync - data is last block received (TCP close after packet sent)
- open continued
- invalid block type (data is 1 byte block type)
- invalid version number (data is 1 byte version)



State Diagram

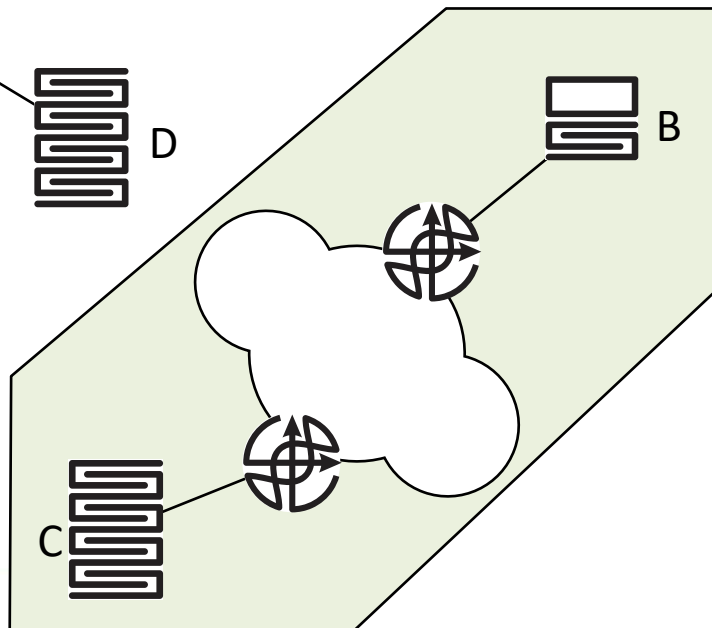
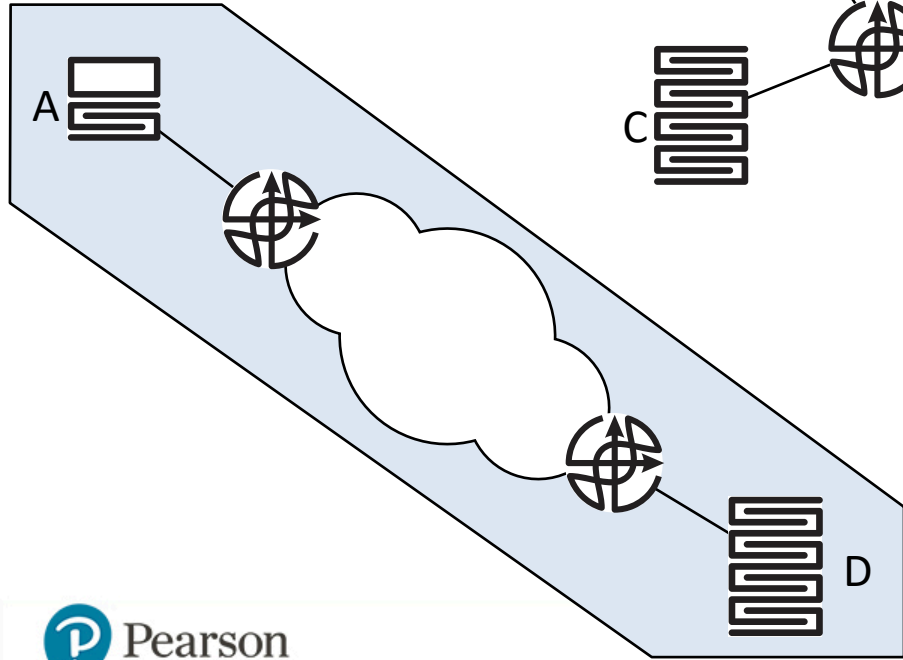
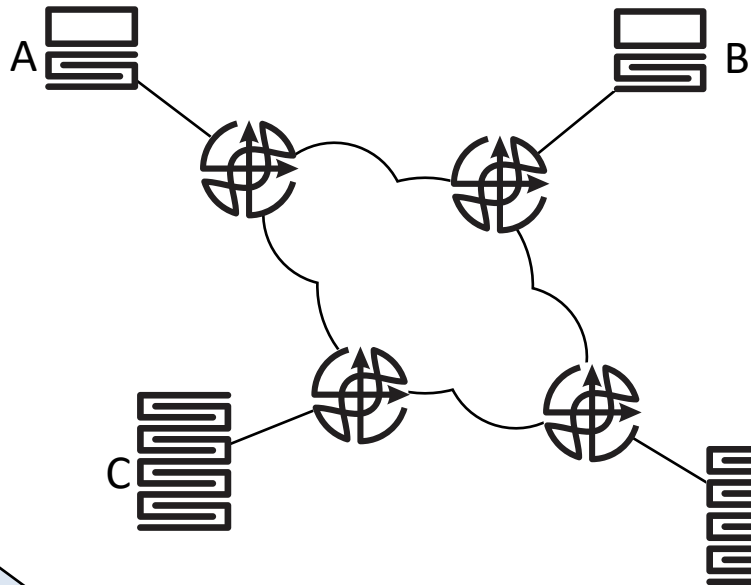


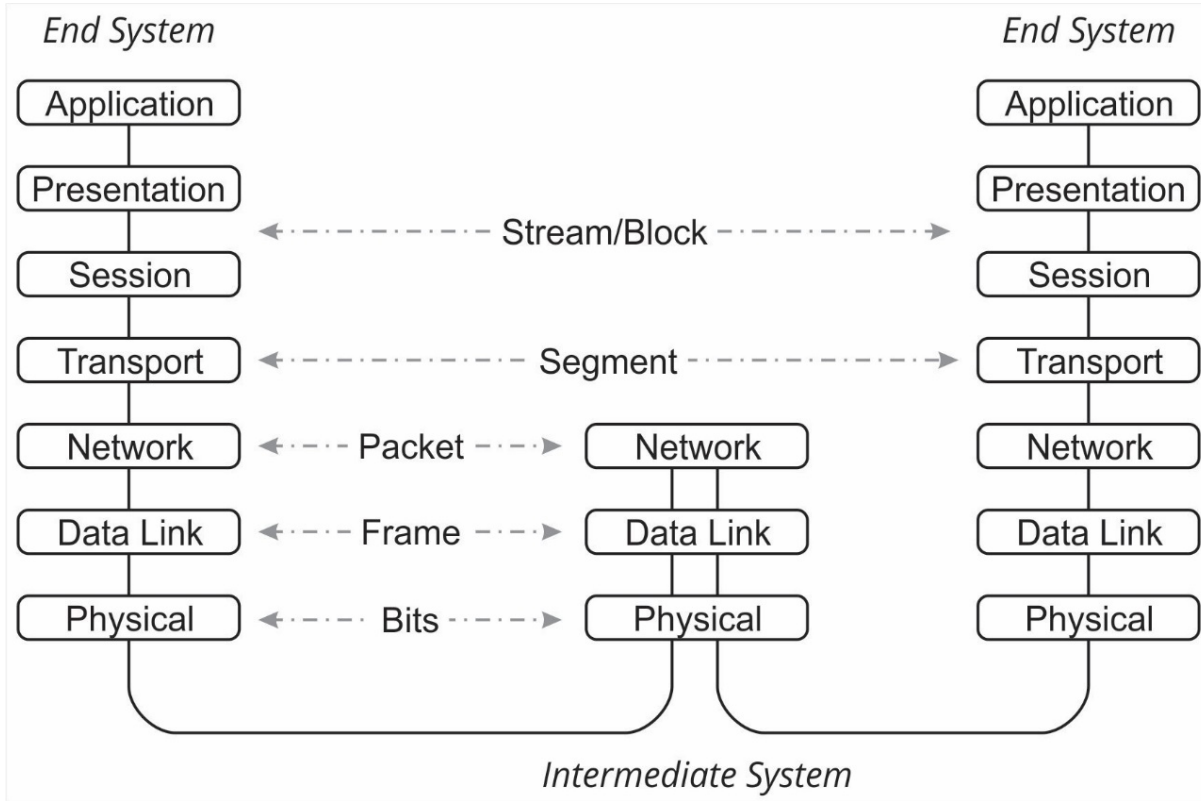
longhead@cisco.com  
YAROV@IBH.COM

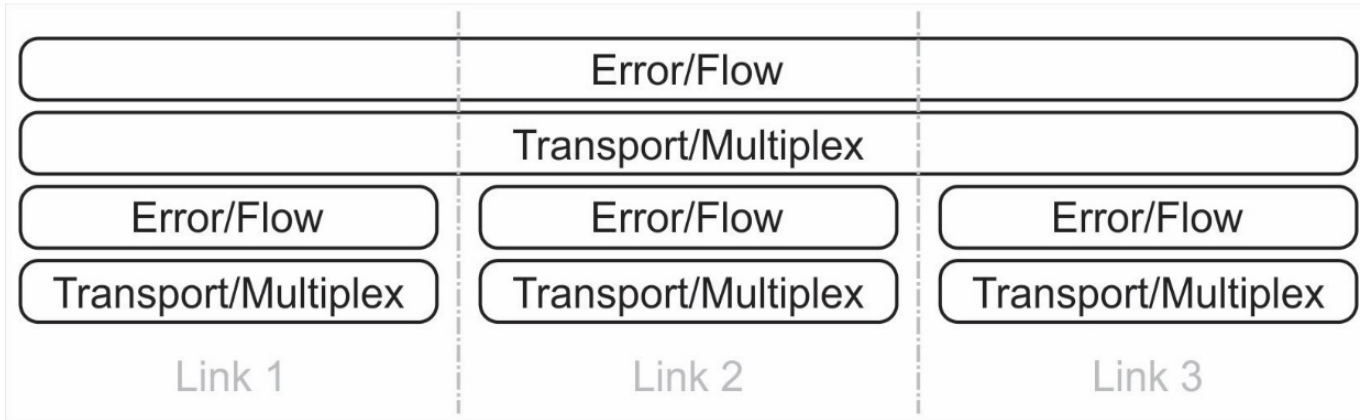
415-326-1941 (11-7) PST  
(914) 415-3896 (8-5) EST













# Measurement Principles

- Measure anything helpful
  - In determining there is a problem
  - In determining where to start troubleshooting
- Do not expect ongoing measurement to provide all the information you need to troubleshoot a problem







# Problems with Observation

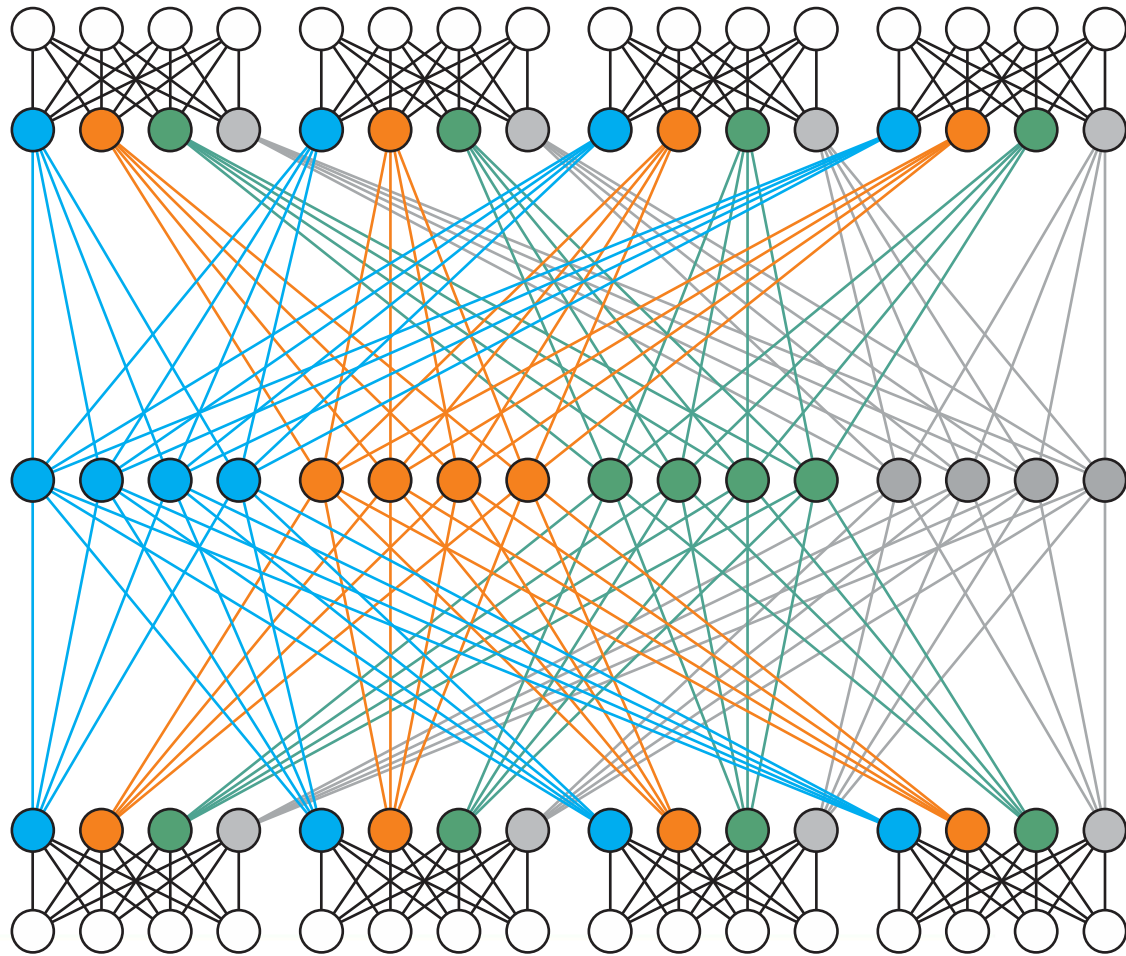
- All these things change the observation
  - Removing links
  - Pinning traffic to a path
  - Showing what is in the queue
  - Just about anything else
- You often cannot observe without changing
- Make certain you take this into account while troubleshooting




# Manipulability Theory

- Prove the problem is at “here” before moving deeper into “here” as a system
- Manipulability theory will help
  - A form of causality theory
  - Closely related to the so-called “scientific method”
  - *Making Things Happen* is a good explanation of this theory, for the more philosophically inclined









Normal Troubleshooting

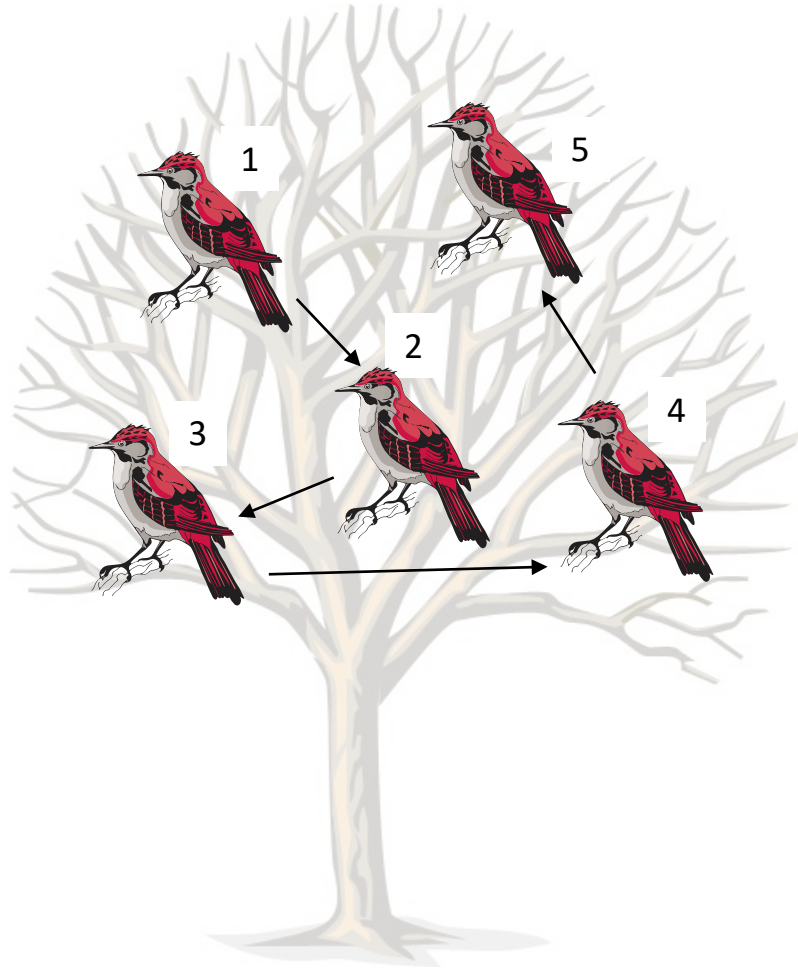
Half Split Method

Some General Notes on Troubleshooting


Three diagnostic heuristics were identified as being in use:


- a) initially look for correlation between the behaviour and any recent changes made in the software,
- b) upon finding no correlation with a software change, widen the search to any potential contributors imagined, and
- c) when choosing a diagnostic direction, reduce it by focusing on the one that most easily comes to mind, either because symptoms match those of a difficult to diagnose event in the past, or those of any recent symptoms match those of a difficult-to-diagnose event in the past, or those of any recent
- d) A fourth heuristic is coordinative in nature: when making changes to software in an effort to mitigate the untoward effects or to resolve the issue completely, rely on peer review of the changes more than automated testing (if at all.)

*Allspaw, John. "Trade-Offs Under Pressure: Heuristics and Observations of Teams Resolving Internet Service Outages." Lund University, 2015.*



=









# Random Walks are Inefficient

- Particularly in the case of troubleshooting
  - You have a scoped space
  - You know something exists that you need to find
  - But it could be something that is missing, rather than something that is there
- And yet...
  - The guided random walk is the most common troubleshooting technique deployed




*“I’ve seen this before,  
and I know the  
solution”*



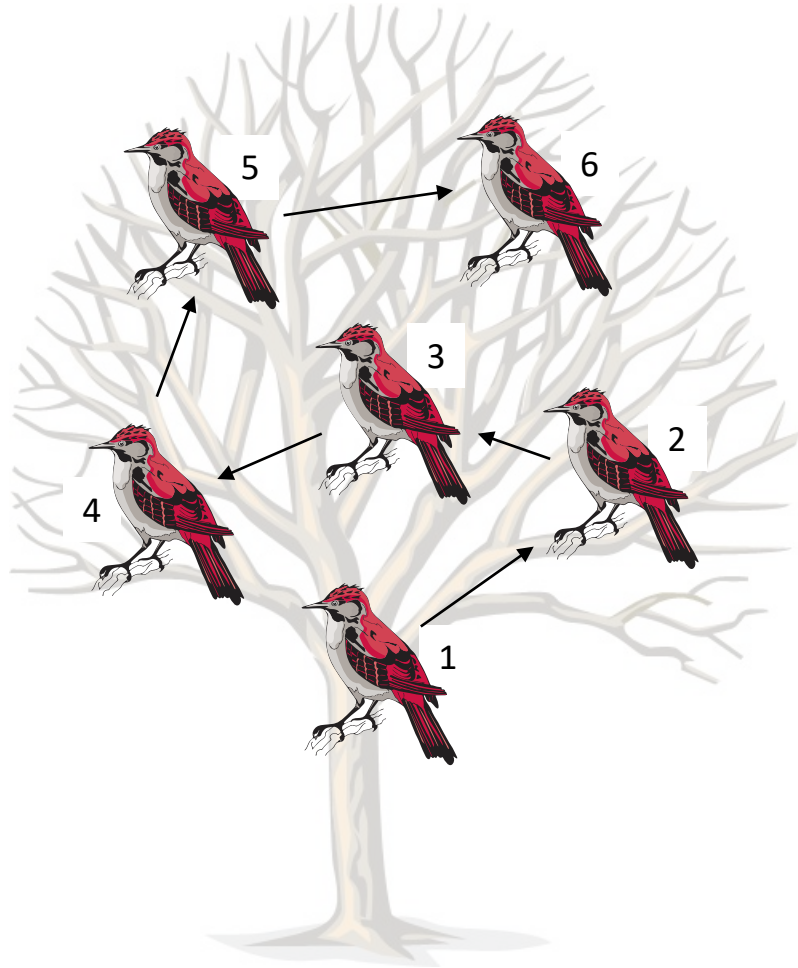
*“I thought I knew the  
problem, but it turns  
out I was wrong”*

**Go back to basics!**

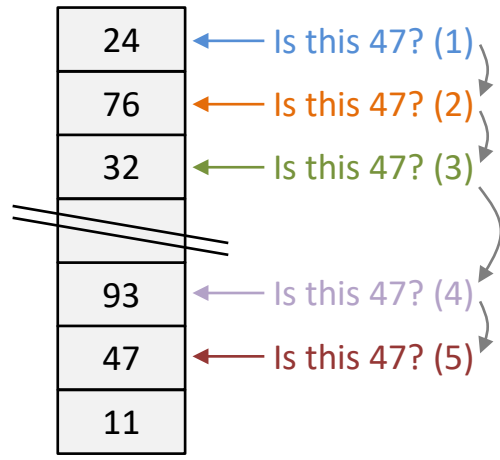



# Avoiding “the Narrows”

- Each of these...
  - If asked before you have narrowed to that point in the network...
  - Will take you down the narrows
- Keep your focus on
  - What the system is supposed to do
  - What the system is doing
- Rather than guessing about what might be wrong
  - Or what you can measure



==





# Linear Process


- More methodical
- Could be useful in some situations
- But not likely to find the problem quickly in most networks...





15, 35, 39, 11, 76, 93, 32, 47, 24




11, 15, 24, 32, 35, 39, 47, 76, 93









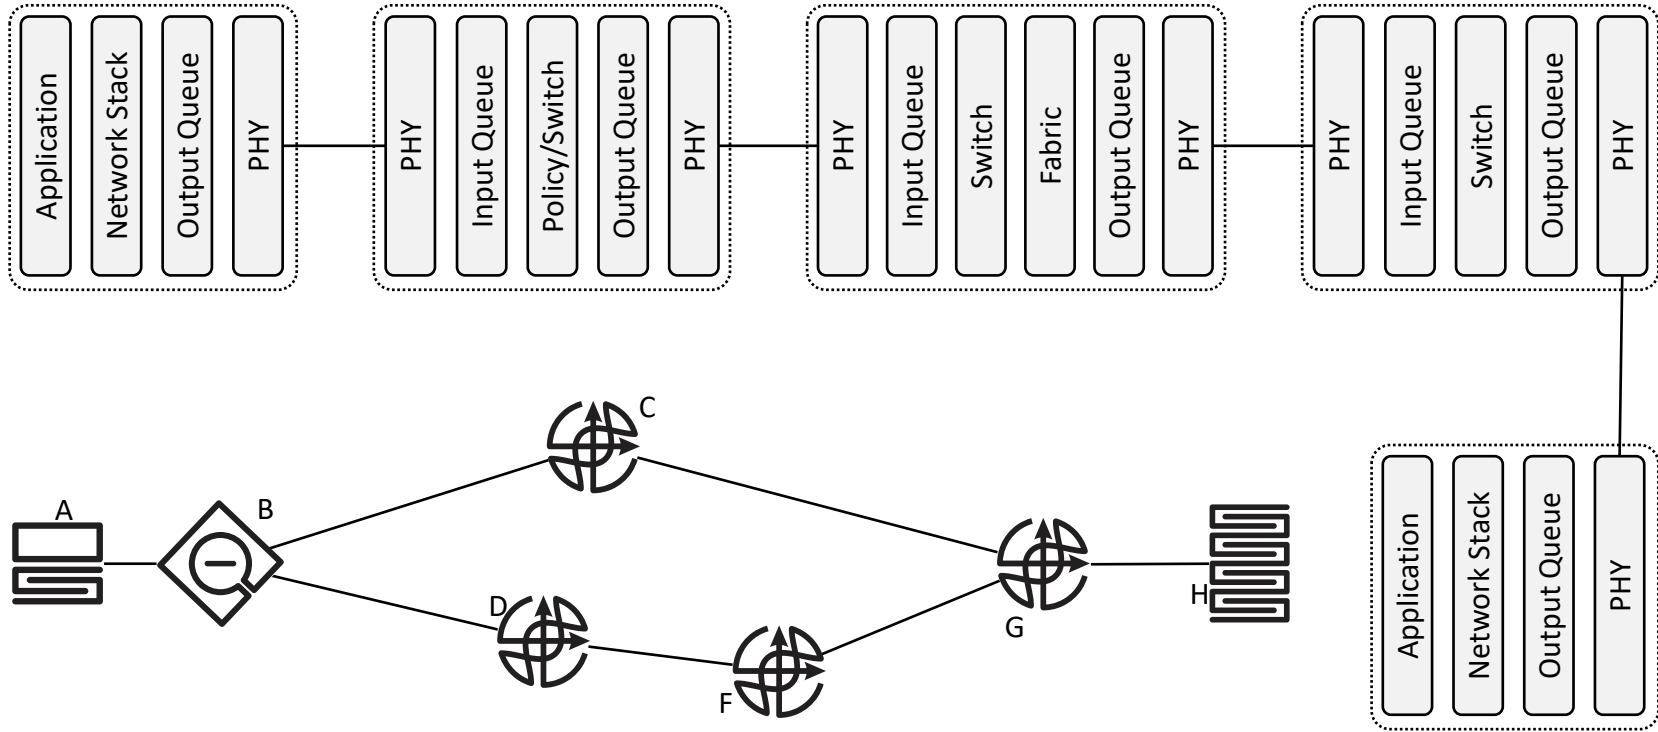




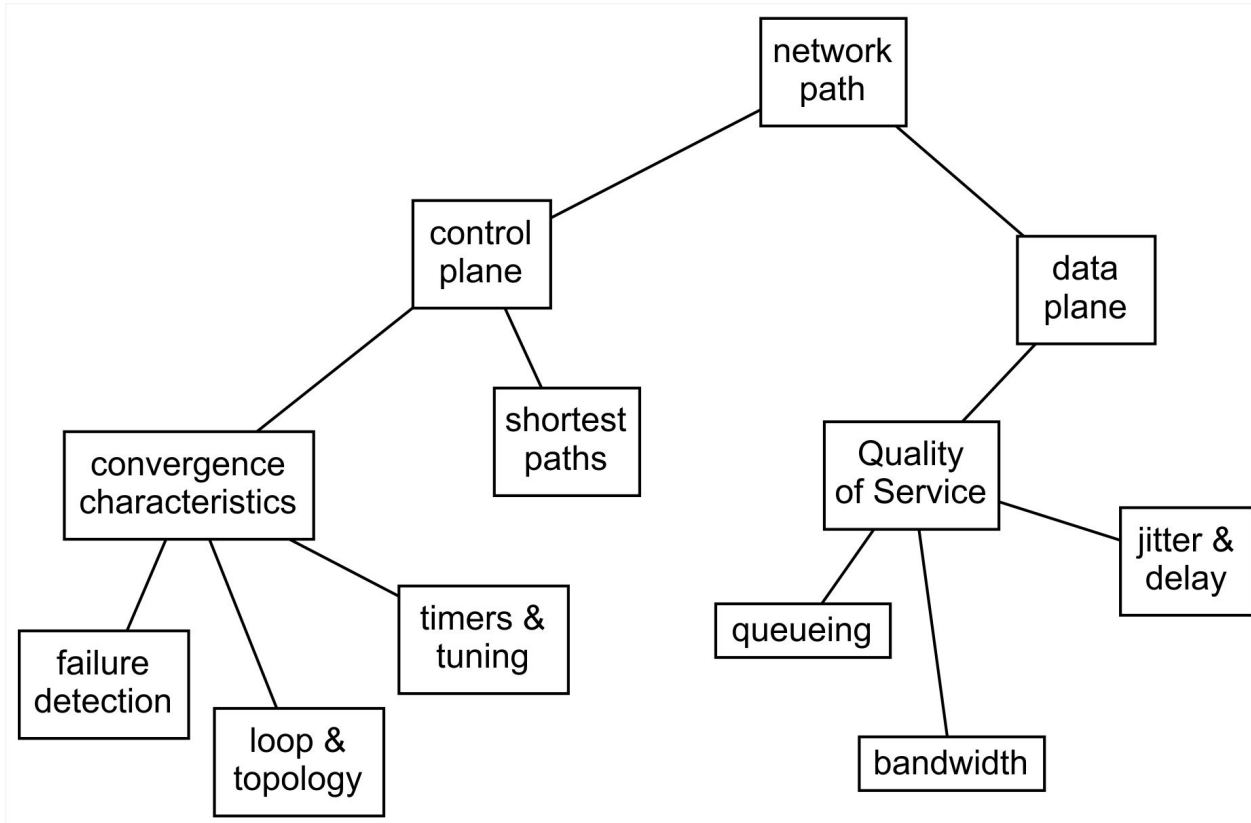
## Putting it All Together: *A Troubleshooting Process*












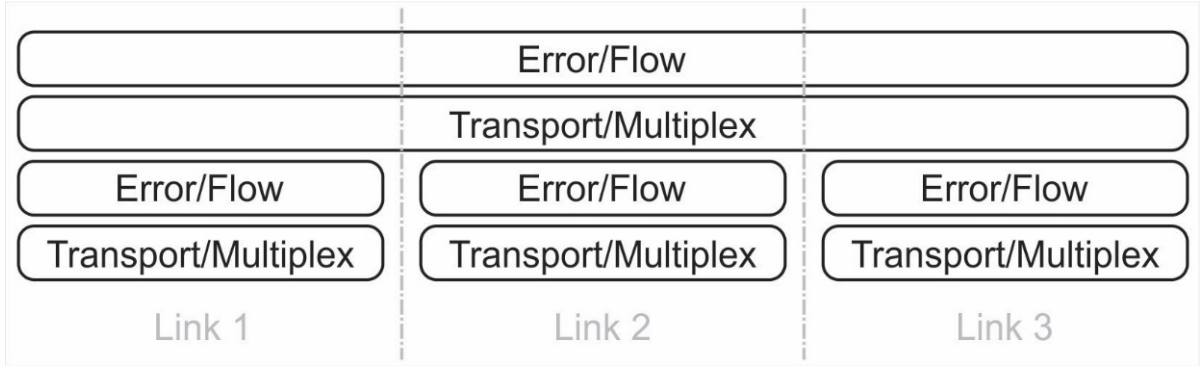
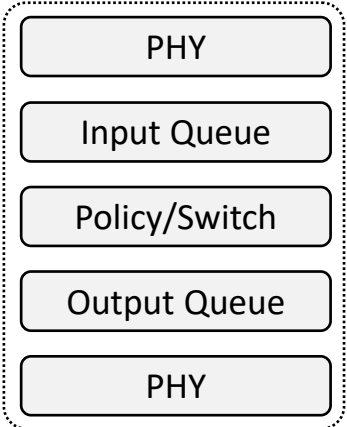
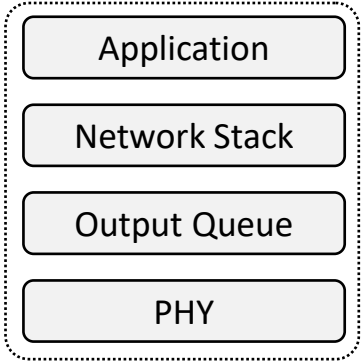


# Discover the Background

- Ask WHAT
  - What do I need to measure?
- To do this I must
  - Understand how this works
  - Understand what I should expect to see
  - Having a solid mental model of the system is going to uncover these things
  - Where are there interfaces between systems?

# Discover the Background

- Ask HOW
  - How can I measure this?
- To do this I must
  - Understand how I can discover whether or not my expectations are correct
  - Understand how I can discover if what I think should be there *is really there*
  - Understand *where* I can measure this effectively





Build lifelines into  
adjacent systems  
*before the failure*

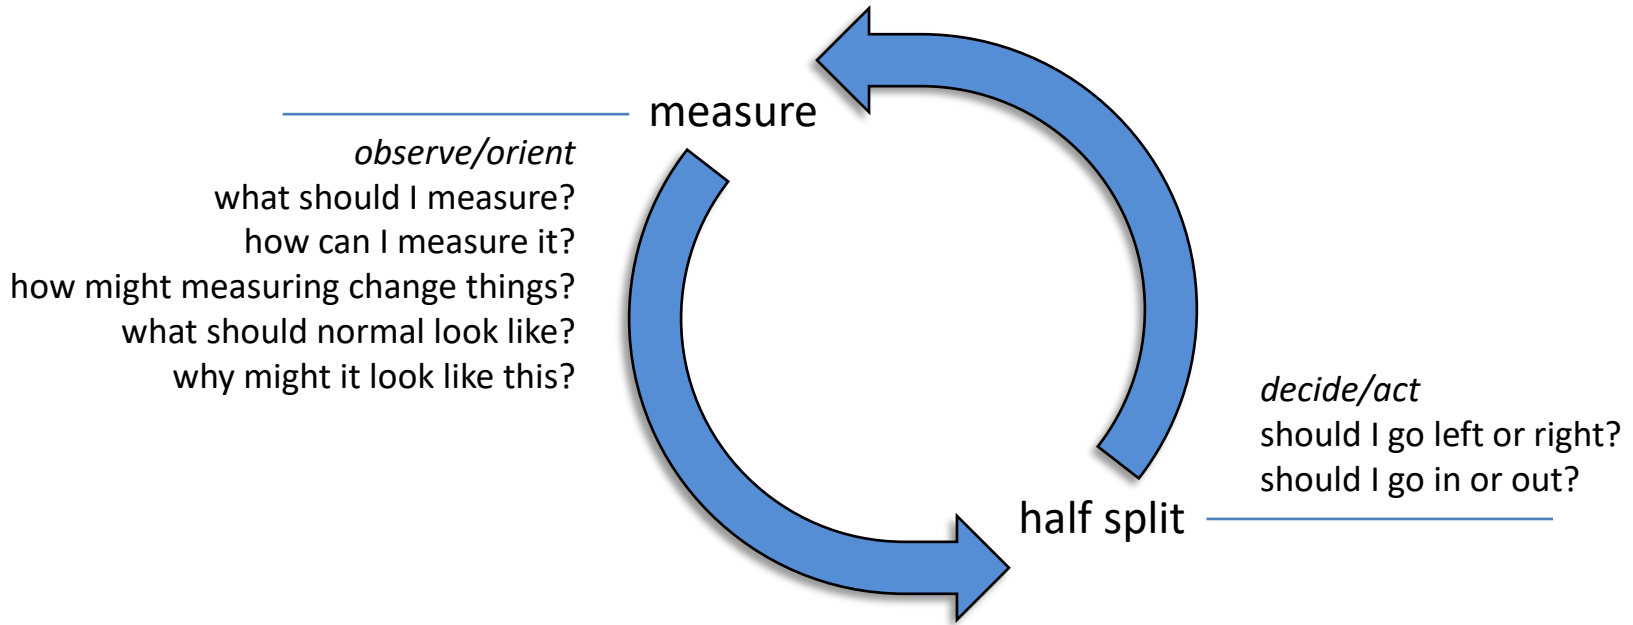
# Summary

- Building understanding before you encounter a failure will help you work faster and avoid the narrows
- Build a baseline of what is normal
  - Gather information from your measurement points periodically
  - Know what things “should” look like
- Model the systems mentally
  - How does this work?
  - Why does this work this way?

# Summary

- Build decision points into the system
- Think about where the half split points are
  - Design your network with choke points to make half splitting easier
- Think about how you can manipulate things
  - To understand what is going on
  - Without causing damage to adjacent systems











# Example

*Flapping EIGRP*













network

router







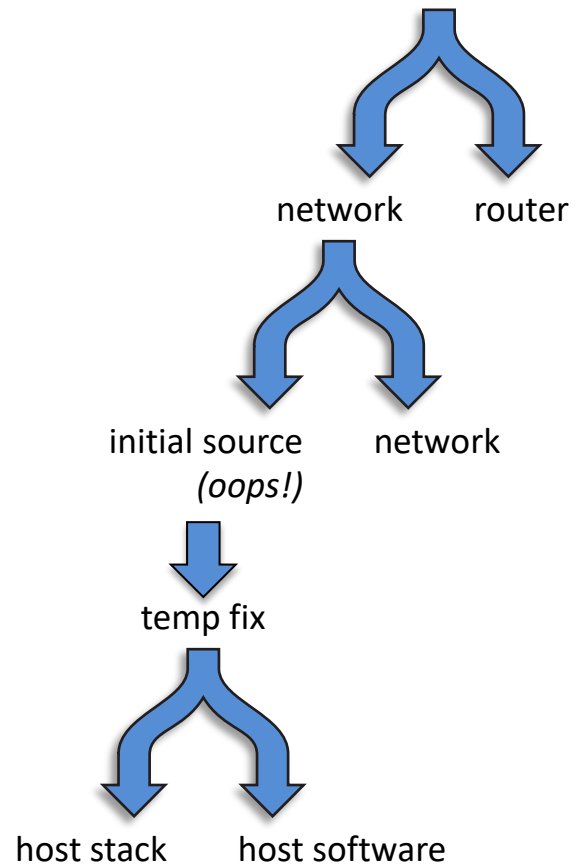


# Trace Information

- Packets are originating from every conceivable source
  - Not just network management hosts
  - So this is not likely a network management application
- We can't measure every host in the network ...
  - Choose a small set to measure



192.0.2.7  
192.0.2.15  
192.0.2.23  
192.0.2.31  
...



# Host Troubleshooting

- What has been installed recently?
  - Ask the obvious question first!
- Measure
  - Disable this software on two hosts machines
  - Leave the software active on two other hosts
  - Note their IP addresses
  - Log IP packets at a core router with these four source addresses
- Result
  - No problematic subnet broadcasts being transmitted by the two modified hosts
- Root cause!


# Root Cause

- Software installed on all machines was intended as the *server half* of a backup package
- To discover clients
  - Sends subnet broadcasts
  - To *every possible* subnet on classful boundaries
- Installed on more than 10k machines
  - Subnet broadcasts being replicated by *every* router in the network
  - Filling input queues
  - EIGRP was just the “canary in the coal mine”




# Example

*Flapping BGP*









bgp configuration


link failures






bgp configuration

link failures



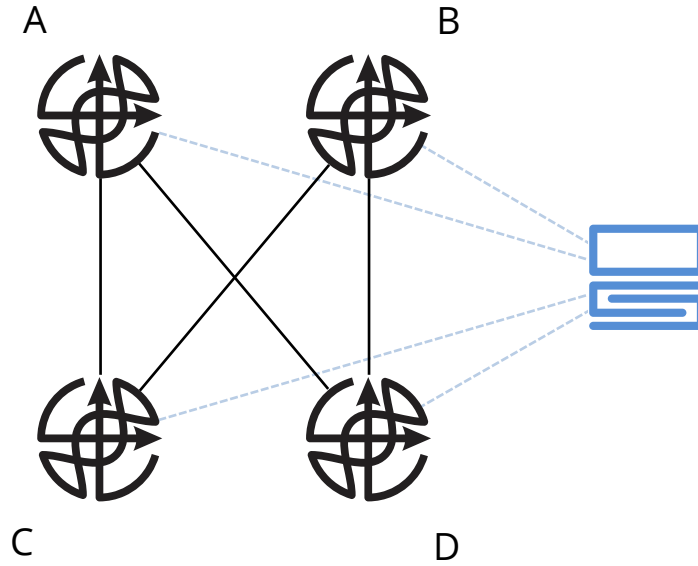



bgp configuration link failures

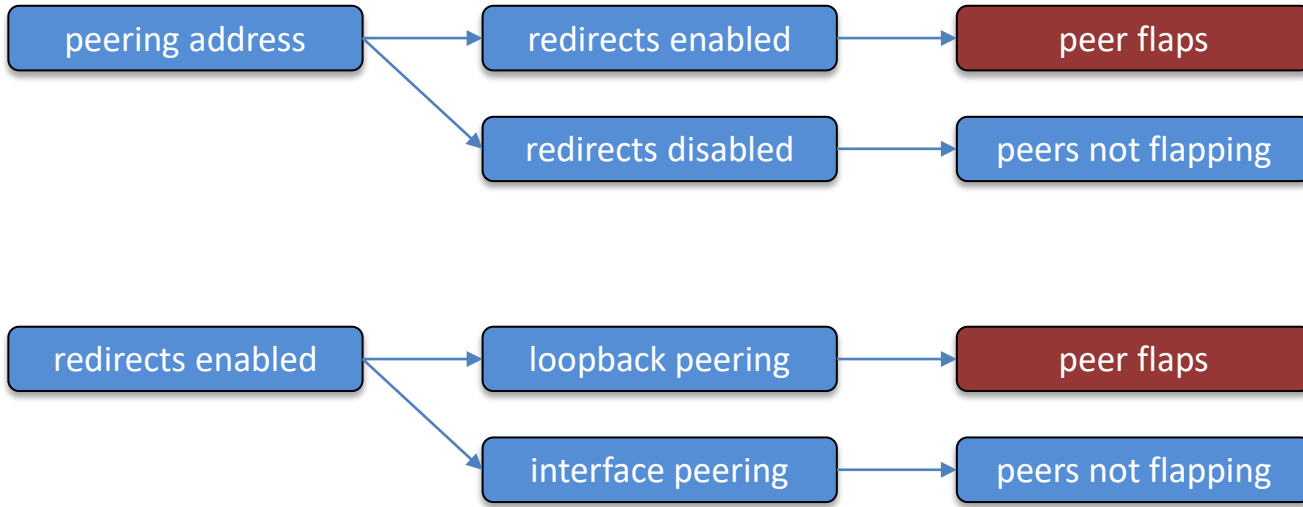


packet drops, etc.

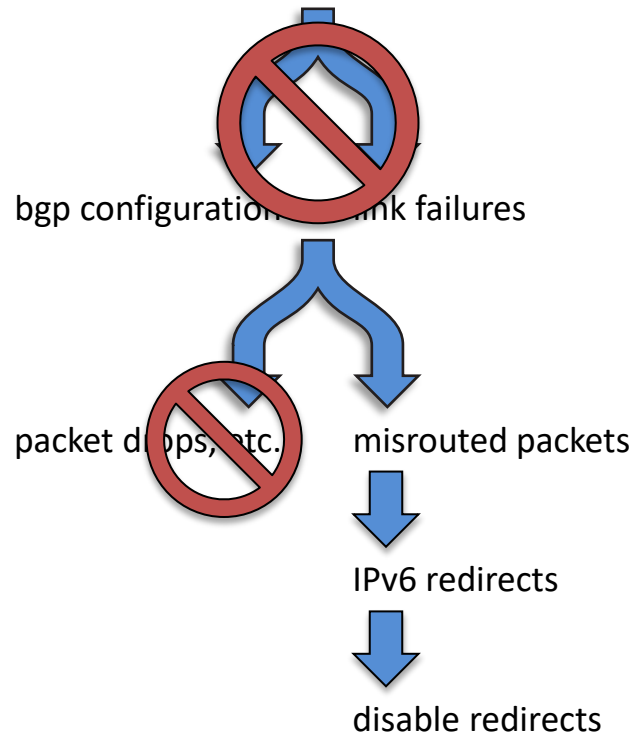
???















# Example

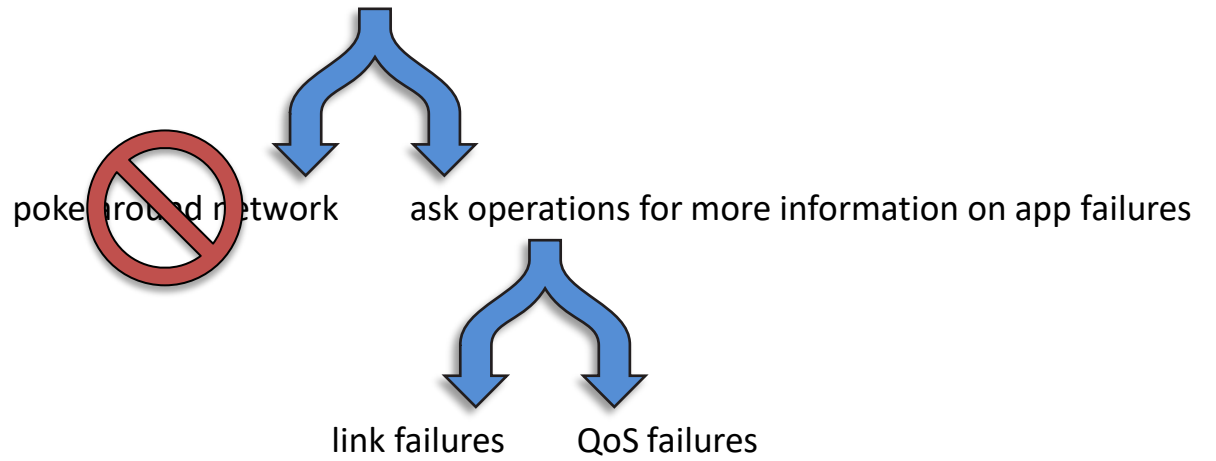
*Application Performance over DCI*







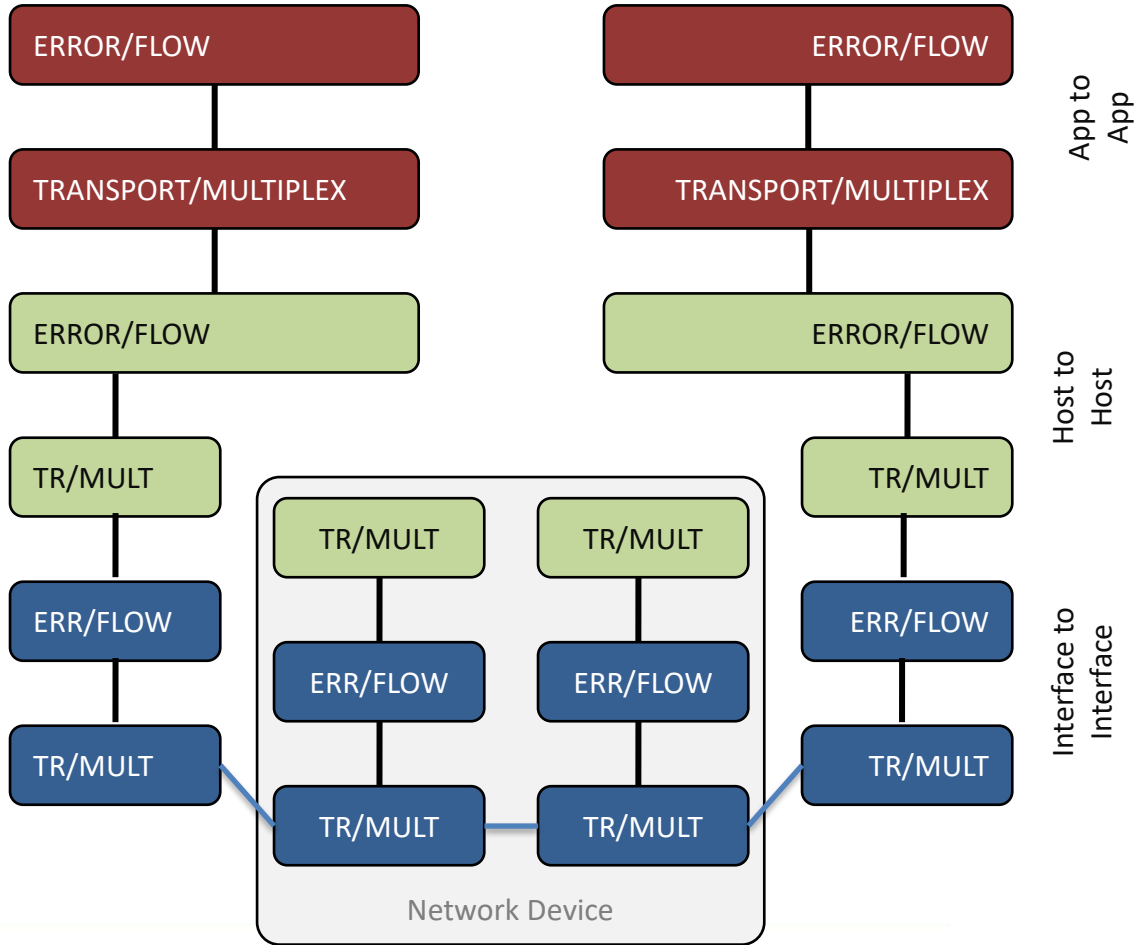
poke around network

ask operations for more information on app failures





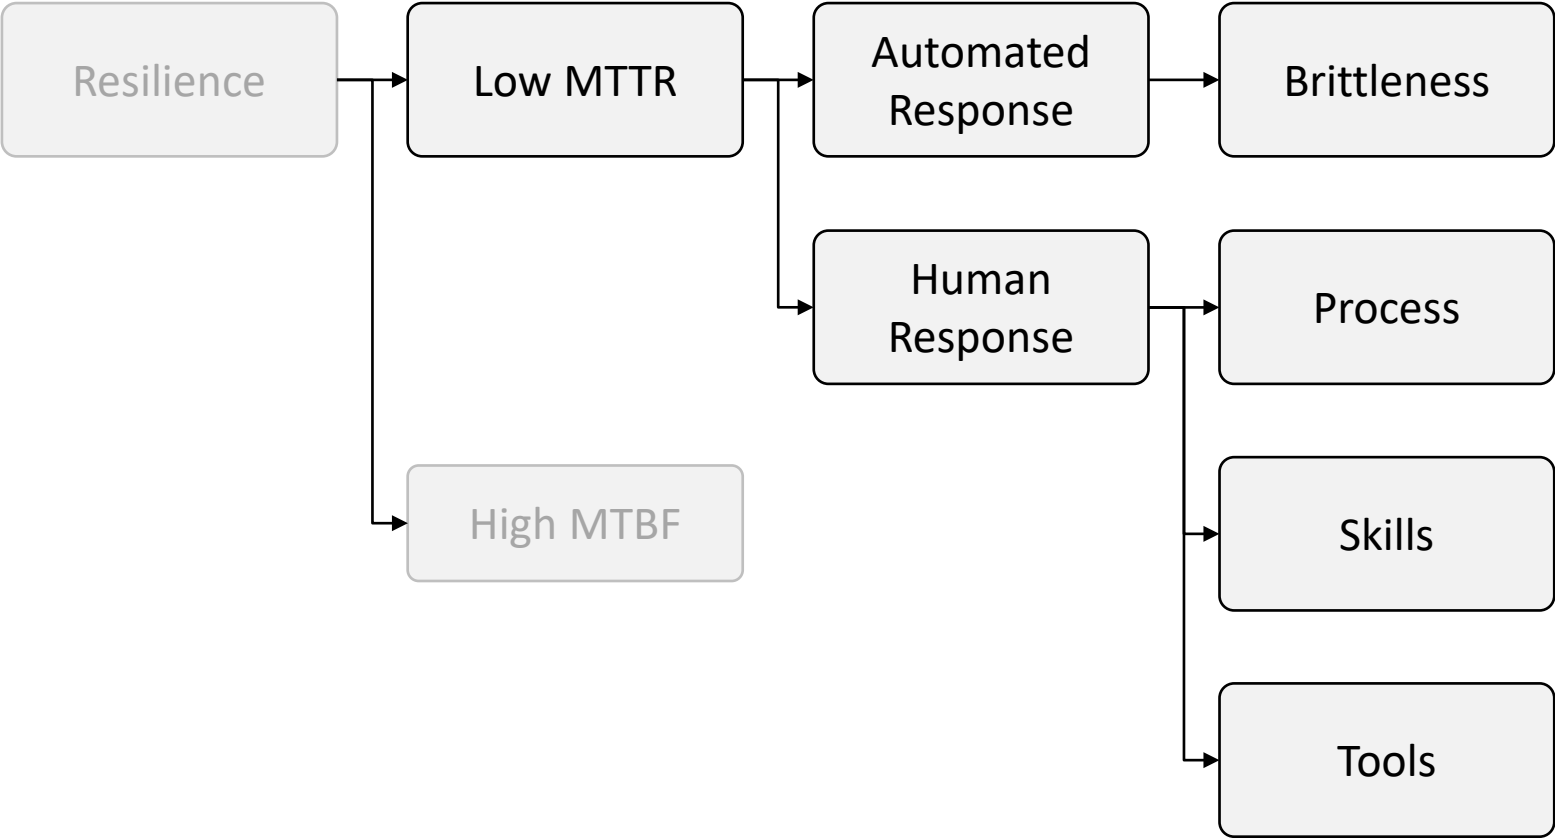









# Summary



# Summary

- Avoid the narrows
  - Record where you've been
  - Be willing to back up if needed
  - Stay focused on the problem you are trying to solve



**AGGREGATION  
AND ITS  
TRADEOFFS**



P Pearson  
*livelesson*

**DATA CENTER  
TOPOLOGIES AND  
CONTROL PLANES**



P Pearson  
*livewebinar*

**THE ART OF  
NETWORK  
ARCHITECTURE**  
BUSINESS DRIVEN DESIGN



P Pearson

**LARGE SCALE  
NETWORK  
DESIGN**



P Pearson  
*livelesson*

**COMPUTER  
NETWORKING  
PROBLEMS AND  
SOLUTIONS**  
AN INNOVATIVE APPROACH  
TO BUILDING RESILIENT, MODERN  
NETWORKS



P Pearson

**DISAGGREGATION  
IN DATA CENTER  
NETWORKS**



P Pearson  
*livelesson*

**INTERMEDIATE  
SYSTEM TO  
INTERMEDIATE  
SYSTEM**



P Pearson  
*livelesson*

**NAVIGATING  
NETWORK  
COMPLEXITY**  
UNDERSTANDING COMPLEXITY  
IN ROUTING AND NEXT  
GENERATION NETWORKS



P Pearson

**HOW  
ROUTERS  
REALLY  
WORK**



P Pearson  
*livewebinar*

